

(a) audio: six speakers in turn




six speech clips, distinct speaker and topic

query names segment k



segment k = one contiguous token run

$+B$ on the target span



added to the pre-softmax attention logits of the K selected heads

(b) vision: a six-panel comic strip



six panels, tiled side by side

query names segment k



panel k = 24 interleaved runs of 7 (stride 42)

$+B$ on the target span



added to the pre-softmax attention logits of the K selected heads